

A Appendix: Hierarchical Code Generation

In this section, we expand upon the source code generation process of Blueprint's toolchain, Blueprint. Recall that the service instance of the `ComposePostService` in the IR was modified with two modifiers - `ZipkinTracer` and `GRPC`. As a result, Blueprint generates two different classes that wrap around the workflow specification from Fig. 1. Fig. 13 shows the generated code. Blueprint generates code in a hierarchical fashion. As tracing must happen before the start of the actual processing of the function call, Blueprint wraps the workflow specification of the service with a generated object that starts a new span every before calling into the function. Subsequently, as the next level in the hierarchy is the `GRPC` node, Blueprint wraps the tracing code with `GRPC`-specific code and creates a function to start the `GRPC` server. The next node in the hierarchy, according to the IR, is the process node. Fig. 14 shows the process code generated for the service. In lines 2-15, the client objects are constructed and then passed as parameters to the constructor for the `ComposePostService` defined in Fig. 1. The tracer and `grpc` objects for the service are then constructed and the server is started. This hierarchical code generation pattern repeats for every service instance defined in the wiring specification of the application.

B Phenomena, Systems Usage, & Modifications

B.1 Known Emergent Phenomena

In this section, we list and define known phenomena. While, our list of phenomena is by no means exhaustive, the list is detailed enough to provide an insight into the kinds of phenomena exhibited by microservice systems and what are the requirements for a system to exhibit phenomena. We hope that researchers find this information useful while using Blueprint to modify systems to understand this phenomena.

Retry storm metastable failure. in which a system is pushed into a prolonged failure state due to repeated RPC retries and timeouts [7, 15, 33]. Retry storms arise when the timeout configuration of services are too low, causing RPCs to be aborted part-way through execution, wasting work. The system must be modified such that each service has timeouts as well as retry mechanisms where a service generates a retry for a request on failure.

Load Spike Resource Contention metastable failure. This is type 2 of the four types of metastability failures [33]. In this phenomenon, some bookkeeping function competes with the actual job for processor resources. For example, a sudden spike in the load could trigger frequent Garbage Collection(GC) cycles which compete with actual job processing for processor resources leading to a build up of unprocessed requests leading to a cascade of backed up requests. For a system to exhibit this phenomenon, the system must have

some kind of an auxiliary bookkeeping subroutine – 100% sampling for traces, garbage collection.

Low-Level Resource Contention metastable failure. This is type 3 of the four types of metastability failures [33]. In this phenomenon, low-level resource contention outside the scope of the application causes the system to slow down. For example, co-location of two processor intensive services might cause those services to drastically impact each other. For a system to exhibit this phenomenon, the system should additionally have timeouts and retries to trigger a cascading effect so that the system stays in a metastable state.

Low-Level Resource Contention QoS violations. Similar to the previous phenomenon, this phenomenon manifests as QoS (Quality of Service) violations caused due to some low-level resource contention. However, in this phenomenon, the violations only last a short duration and the system does not enter into a metastable state. For a system to exhibit this phenomenon, the system needs some kind of low-level resource contention that violates a QoS metric. An example would be FIRM's anomaly injector [55] that injects low level anomalies to cause QoS violations.

Capacity Degradation Trigger Capacity Degradation Amplification metastable failure, This is type 4 of the four types of metastability failures [33]. In this phenomenon, a trigger decreases the capacity of the system to process requests which causes failures leading to the capacity of the system never reverting back to the normal state. For example, a system with a cache in front of a database can process 2 Kbps but if the cache fails then the capacity of the system to process requests decreases. This can cause failures and prevent the cache from ever getting filled back and can overload the database. For a system to exhibit this phenomena, the system needs two different paths - a fast path and a slow path and timeouts and retries. Momentary failure of the fast path could lead to timeouts causing a retry storm and putting the system into a metastable failure state.

Cascading QoS violations in which application quality of service deteriorates due to increased latency and/or failures in downstream services [52, 75]. This is part of a broader class of *cascading effects* where failures in one component drastically combine with error-handling and availability mechanisms in another component to ultimately bring down the system [70]. For a system to exhibit this phenomenon, there must be a chain of services of significant length as well as timeouts and retries.

QoS violations from repeated expensive I/O. In this phenomenon, repeated execution of expensive I/O can lead to QoS violations. For example, repeatedly fetching the same information from a resource that is expensive to access, in terms of I/O overhead or latency and not caching the data can cause QoS violations [9]. Another example would be services not re-using clients while making requests to other

```

1 type ComposePostTracer struct {
2     service *ComposePostImpl
3     tracer blueprint.Tracer
4     service_name string
5 }
6
7 func NewComposePostTracer(service *ComposePostImpl, tracer blueprint.
8     Tracer, service_name string) *ComposePostTracer {
9     return &ComposePostTracer{service: service, tracer: tracer,
10         service_name: service_name}
11 }
12
13 func (t *ComposePostTracer) ComposePost(ctx context.Context, reqID
14     int64, username string, userID int64, text string, mediaIDs []
15     int64, mediaTypes []string, post_type services.PostType,
16     zipkinTracer trace.Context) error {
17     if zipkinTracer.trace_ctx != "" {
18         span_ctx_config := blueprint.GetSpanContext(
19             zipkinTracer.trace_ctx)
20         span_ctx := trace.NewSpanContext(span_ctx_config)
21         ctx = trace.ContextWithRemoteSpanContext(ctx, span_ctx)
22     }
23     tp, _ := t.tracer.GetTracerProvider()
24     tr := tp.Tracer(t.service_name)
25     ctx, span := tr.Start(ctx, "ComposePost")
26     defer span.End()
27     err := t.service.ComposePost(ctx, reqID, username, userID, text,
28         mediaIDs, mediaTypes, post_type)
29     if err != nil {
30         span.RecordError(err)
31     }
32     return err
33 }

```

(a) Generated code for adding tracing to ComposePostService

```

1 type ComposePostGRPC struct {
2     service *ComposePostTracer
3     socialnetwork.UnimplementedComposePostServiceImplServer
4 }
5
6 func NewComposePostGRPC(old_handler *ComposePostServiceImpl) error {
7     addr := os.Getenv("composePostService_ADDRESS")
8     port := os.Getenv("composePostService_PORT")
9     if addr == "" || port == "" {
10         return errors.New("Address or Port were not set")
11     }
12     lis, err := net.Listen("tcp", addr + ":" + port)
13     if err != nil {
14         return err
15     }
16     handler := &ComposePostGRPC{service: old_handler}
17     grpcServer := grpc.NewServer()
18     socialnetwork.RegisterComposePostServiceImplServer(grpcServer,
19         handler)
20     return grpcServer.Serve(lis)
21 }
22
23 func (rpchandler *ComposePostGRPC) ComposePost(ctx context.Context,
24     request *socialnetwork.
25     ComposePostServiceImpl_ComposePostRequest) (*socialnetwork.
26     BaseRPCResponse, error) {
27     var arg7 services.PostType
28     copier.Copy(&arg7, request.PostType)
29     ret0 := rpchandler.service.ComposePost(ctx, request.ReqID, request.
30         Username, request.UserID, request.Text, request.MediaIDs, request.
31         MediaTypes, arg7, request.ZipkinTraceCtx)
32     response := &socialnetwork.BaseRPCResponse{}
33     return response, ret0
34 }

```

(b) Generated code for running ComposePostService as a GRPC server

Fig. 13. Generated Server side code for the ComposePostService defined in §4

```

1 func RuncomposePostService() error {
2     zipkin := NewZipkinTracer()
3     userserviceimplrpcclient_netclient, err :=
4         NewUserServiceGRPCClient()
5     if err != nil {
6         return err
7     }
8     userservicetracerclient := NewUserServiceZipkinClient(zipkin,
9         userserviceimplrpcclient_netclient)
10    userserviceimplclient := NewUserServiceClient(
11        userservicetracerclient)
12
13    poststorageserviceimplrpcclient_netclient, err :=
14        NewPostStorageServiceGRPCClient()
15    if err != nil {
16        return err
17    }
18    poststorageserviceimpltracer :=
19        NewPostStorageServiceZipkinClient(zipkin,
20        poststorageserviceimplrpcclient_netclient)
21    poststorageserviceimplclient := NewPostStorageServiceClient(
22        poststorageserviceimpltracer)
23    spec_handler := NewComposePostImpl(poststorageserviceimplclient,
24        userserviceimplclient)
25    composeposttracer := NewComposePostZipkin(spec_handler)
26    err_new := NewComposePostServiceImplHandler(composeposttracer)
27    if err_new != nil {
28        return err_new
29    }
30    return nil
31 }

```

Fig. 14. Generated Server side main function for the ComposePostService defined in §4

services and re-establishing a connection on every call. For a system to exhibit this phenomenon, the system must perform

an expensive operation in a repeated fashion either in the workflow spec or in the scaffolding.

Reduced Availability due to saturation. Services can get saturated and overwhelmed when placed under significant load. However, even in the absence of retry storms, a service can get oversaturated. In this phenomenon, only one service is impacted while the rest of the system remains healthy. This can happen in multiple scenarios. One such scenario is if the service has a high fan in number *i.e.* multiple services call into this service. Another scenario is when one big request is broken down into multiple smaller requests which can fill up queues at the called service [8]. DSB-SN v1 suffered from this exact phenomenon via a combination of the above scenarios where multiple services called ComposePostService and the items to ComposePostCache were broken down into multiple pieces and accessed via multiple set/get requests. The application was changed to remove this phenomenon [19].

Reduced Availability due to QoS violation prevention schemes. In this phenomenon, application availability is impacted by mechanisms such as circuit breakers [60] in place to prevent QoS violations at others services. This incurs a high error rate or high end-to-end latencies to prevent QoS violations at others services. For a system to exhibit this phenomenon, the system must have QoS prevention mechanisms in different parts of the system.

Cross-System Inconsistency. This phenomenon occurs when requests write to multiple eventually-consistent datastores, establishing an ordering for read operations across those datastores. Each datastore is individually consistent, yet readers may experience inconsistent reads, such as a notification arriving before post content has finished replicating [3]. For a system to exhibit this phenomenon, the system must at least have replicated datastores. For a stronger manifestation of this phenomenon, geo-replication is recommended [42].

User-View Data Inconsistency. This phenomenon is a less stricter version of Cross-System Inconsistency. For a Cross-System Inconsistency, the replicated datastore is inconsistent from the viewpoint of the system itself such that a value that was committed in the datastore has not propagated to all replicas. In contrast, for a user-view data inconsistency, the value does not have to be committed to a datastore; rather, it must *seem* to be committed from the user but the update has not reached all parts of the system due to asynchrony [41]. Prime example of this phenomenon is observed almost on a daily basis by users of social-media applications where the user gets a notification but clicking on the notification does not show the update simply because the update is being applied asynchronously as the data policy is of eventual consistency. The key to replicating this phenomenon is to have asynchronous updates to datastores leading to inconsistent view of the system from the point of view of the user. This phenomenon falls under the larger class of phenomena related to Data Integrity violations [14].

QoS violations from excessive repeated calls. In this phenomenon, repeatedly contacting a datastore or a service for the same piece of information can overwhelm the downstream service resulting into a cascading failure. For a system to exhibit this phenomenon, it must make repeated calls for the same data and must not use a cache [9].

QoS violations from excessive data demand. This is the dual of the previous phenomenon. In this phenomenon, contacting a service for more data than required can cause the downstream service to get blocked as it now needs to process and return a large amount of data. This can also create contention on the network and bog down concurrent requests. For a system to exhibit this phenomenon, the workflow spec must have a service that asks for excessive data.

Performance Degradation due to Disruptive neighbours. In this phenomenon, the performance of one application/tenant is drastically impacted by the co-located application/tenant due to the co-tenant using a disproportionate amount of resources [10, 45]. This scenario only happens in situations when multiple applications are co-located on the same node and compete for resources.

The Laser of Death. In this phenomenon, the availability of the system is hampered by a load balancer overloading replicas of a service [51]. Specifically, a load balancer performs periodic health checks to find a healthy subset of replicas and routes requests to this healthy subset. If one of the replicas gets overloaded then its health check would fail which would cause the load balancer to stop routing requests to that replica. This in turn can overload the new subset of healthy replicas as each healthy replica now has to handle an increasing workload. Essentially, the load balancer begins acting as a Laser of Death by overloading healthy replicas due to health checks. This scenario only happens in applications with replicated service instances that are fronted by a load balancer making routing decisions based on health checks of instances.

The Killer Health Check. In this phenomenon, health check based instance replacement automation might cause warmed up hosts to get replaced which can prevent the application from servicing requests at the same rate [51]. The underlying cause of this is the fact that the health check signal does not actually specify if a problem is instance-specific. The problem might exist in some downstream services or might be a problem that is affecting all services. This scenario can only happen in applications with replicated service instances that allow for automatic creation and deletion of instances based on health check signals.

Incorrect Microservice Granularity. This is a phenomenon pertinent exclusively to the hierarchy and grouping of microservices. Mega Microservices can suffer the same problems that monolith applications suffer whereas Nano Microservices can incur a high development and maintenance and can lead to cycles [69]. The correct granularity of microservices has also been of interest to practitioners. This interest led to Uber switching to a less granular microservice architecture in which microservices are grouped together into domains [167].

Multi-version Errors. This is a phenomenon in which multiple versions of the same microservice co-exist and are deployed at the same time [69]. This can cause a lot of errors if the requests meant for service:v1 end up at service:v2. For a multi-version error to occur, the system could have service discovery mechanisms that link callers to incorrect version due to stale information. Another way of exhibiting this phenomena in this system is to have the same service instance simultaneously service requests from both versions with clients making calls to version 1 with data for version 2 of the function.

B.2 Existing Use of Microservice Systems

To understand how microservice systems are used by researchers, we conducted a literature survey to analyze which systems were popular. To find the necessary papers, we

System	Count	Papers
DSB-SN [26]	51	[34, 40, 52, 55, 82, 84, 85, 95, 102, 103, 107, 108, 126, 129, 133, 149, 156, 157, 160, 162, 191, 196, 201–203, 206, 209, 228, 252, 270, 273, 276–278, 285–287, 304, 322, 323, 334, 351, 353, 371, 374–377]
DSB-MM [26]	18	[55, 103, 107, 108, 111, 131, 133, 156, 162, 209, 225, 252, 273, 286, 287, 305, 322, 337]
DSB-HR [26]	17	[55, 103, 126, 156, 162, 188–190, 205, 209, 217, 246, 302, 314, 375, 384, 388]
DSB-Swarm [26]	1	[163]
μ Suite [66]	2	[253, 328]
TrainTicket [76]	52	[34, 55, 77, 81, 88, 95, 109, 110, 113, 127, 130, 131, 134, 137, 156, 177, 182, 188–192, 197, 203, 234, 235, 237, 239, 246, 247, 257, 296, 297, 304, 305, 311, 322, 329, 338, 339, 347–349, 351, 370, 378–381]
SockShop [47]	96	[83, 86, 87, 91, 93, 100, 105, 106, 114, 115, 117, 119, 121, 124, 134, 135, 139, 141, 146, 150, 155, 171, 172, 176, 178–183, 188–190, 204, 207–211, 215, 216, 218–222, 229–232, 237, 244, 246, 247, 249, 250, 254, 256, 258–262, 268, 271, 280, 281, 283, 286, 292, 294–296, 300, 301, 311, 313, 319–321, 331–333, 335, 336, 340, 343, 350, 354–357, 360–362, 365, 372, 380, 383]
AcmeAir [2]	26	[79, 96, 97, 104, 125, 136, 148, 173, 174, 195, 196, 204, 214, 255, 269, 289, 293, 310, 311, 330, 341, 342, 358, 363, 364]
TeaStore [71]	32	[92, 112, 129, 132, 143, 145–147, 154, 166, 168, 175–177, 184, 185, 187, 194, 198, 198, 213, 233, 248, 286, 303, 306, 307, 317, 318, 326, 344, 359]
HipsterShop/OnlineBoutique [28]	34	[52, 94, 95, 101, 116, 118, 120, 122, 123, 158, 159, 170, 186, 212, 217, 227, 231, 232, 237, 245, 250, 267, 290, 296–299, 302, 321, 366–369, 384]
AliBaba [44]	2	[242, 373]
Stan’s RobotShop [72]	9	[52, 98, 123, 206, 236, 259, 266–268]
BookInfo [36]	12	[52, 80, 90, 98, 123, 231, 232, 236, 236, 243, 263, 296, 312]
eShopOnContainer [140]	13	[153, 164, 165, 200, 223, 247, 260, 265–268, 315, 324, 346]
LakeSideMutual [275]	10	[41, 64, 144, 164, 165, 266, 282, 315, 316, 385–387]
eShoppers [308]	3	[127, 128, 339]
PitStop [142]	2	[232, 264]
Lab Insurances Portal [6]	2	[206, 240]
Overleaf [272]	1	[206]
Retwis [288]	3	[163, 217, 309]

Tab. 6. Microservice systems used in the literature.

started with a list of seed microservice systems (DeathStar-Bench [26] and TrainTicket [76]) and shortlisted the papers citing the seed systems that we found using Google Scholar. During the analysis of the shortlisted papers, if we found a previously unseen system used by the paper then we modified our list of systems to include the newly found system. When we added a new system, we increased the shortlisted papers with the papers citing the new system. Some of the systems did not have an accompanying research paper so it was not straightforward to find the papers using those systems via Google Scholar. To overcome this, we used the url of the system as a query to Google Scholar search to find candidate papers.

To decide if a paper used a particular system, we analyzed the paper’s evaluation section. If the paper performed at least one experiment using a given system, the paper was deemed to have used the system. In total, we found 290 research papers that altogether used 20 different systems. Six of the twenty systems, DSB-SN, DSB-HR, DSB-Swarm, DSB-MM, TrainTicket, and TeaStore, were described by the authors of the systems as microservice benchmarks, eleven systems were demonstration apps by a specific framework or a library to showcase some feature of their library or feature in the context of microservice systems, two systems were constructed based on public trace datasets, and the final system,

Overleaf [272] was the sole microservice system that is currently deployed and used publicly. While our survey found 20 systems that are used by researchers in the literature, there exist a lot more open source microservice systems [284].

Tab. 6 shows the breakdown of papers found in the survey categorized by the open-source system they used. Out of all the used systems, SockShop [47] was used by the most number of papers despite its small number of services. We found it surprising that SockShop was used in more papers than any of the systems labeled as ‘microservice benchmarks’. We believe that this is probably due to the fact that the system has been around for a longer period of time as compared to the benchmarks and is enabled with a plethora of orchestration and monitoring features that are desirable to researchers. This indicates that researchers want systems that have a variety of features.

While researchers use microservice systems to analyze and solve specific emergent phenomena, other use-cases of microservices include using them as sources of log and trace data, or as individual datapoints in a larger dataset of microservice systems for source code analysis. The use-cases of microservice systems are diverse and require a diverse set of systems that vary across a variety of dimensions. Thus, it is impossible that any single implementation of a microservice system can be used as the ideal benchmark system for all possible research scenarios.

System	Total Forks	Forks w/ Modifications
DeathStarBench	240	85
TrainTicket	130	25
TeaStore	96	36

Tab. 7. Modified Forks for each Microservice Benchmark

of DeathStarBench that modified the system to induce a metastable failure [22, 152] into the system and another fork that modified the DeathStarBench system to include geo-replication to analyze cross-system inconsistencies [42]. Modifications were not restricted to emergent phenomena. Some other forks modified the systems to understand the impact of various of design choices- for eg, thrift vs grpc [99], energy efficiency of monitoring tools [20], among others. We also found a fork that modified TeaStore to study, reproduce, and exploit CVE-2021-44228 [291]. Researchers have also modified systems to support their experiments [5, 20, 74].

In conclusion, researchers modify systems for one main reason that can manifest as different types of modifications. We found that the reason for modifications was to include some feature that is absent from the system but is necessary for research.

B.3 Modifications to open-source systems

In addition to the literature survey, we also analyzed forks of popular microservice systems to figure out what modifications were being made to the systems. We chose to analyze the forks of the systems labeled as microservice benchmarks as these systems were designed as microservice benchmarks and are the ideal targets for researchers. We analyzed 240 forks of DeathStarBench, 130 forks of TrainTicket, and 96 forks of TrainTicket. Out of the total 464 forks, 146 forks made modifications to the benchmark systems. Tab. 7 shows the breakdown of the modified forks by system.

We found multiple different types of modifications made to the systems. The most common modification to systems were to either enable or disable tracing and other monitoring tools. This included switching on/off tracing, metrics, and logging. Another common modification was to generate the manifest files for deploying systems on a specific framework such as Kubernetes or OpenShift. There were multiple forks that modified systems such that they exhibited a specific emergent phenomena. We found one fork

Appendix References

- [77] A. S. Abdelfattah. Student research abstract: Microservices-based systems visualization. *arXiv preprint arXiv:2207.12998*, 2022.
- [78] AcmeAir. Acmeair. <https://github.com/acmeair>.
- [79] S. Agarwal, R. Sinha, G. Sridhara, P. Das, U. Desai, S. Tamilselvam, A. Singhee, and H. Nakamuro. Monolith to microservice candidates using business functionality inference. In *2021 IEEE International Conference on Web Services (ICWS)*, pages 758–763. IEEE, 2021.
- [80] S. Allen, M. Toslali, S. Parthasarathy, F. Oliveira, and A. K. Coskun. Tritium: A cross-layer analytics system for enhancing microservice rollouts in the cloud. In *Proceedings of the Seventh International Workshop on Container Technologies and Container Clouds*, pages 19–24, 2021.
- [81] N. Alshuqayran. *Static Microservice Architecture Recovery Using Model-Driven Engineering*. PhD thesis, University of Brighton, 2020.
- [82] M. Amaral, T. Chiba, S. Trent, T. Yoshimura, and S. Choochotkaew. MicroLens: A performance analysis framework for microservices using hidden metrics with bpf. In *2022 IEEE 15th International Conference on Cloud Computing (CLOUD)*, pages 230–240. IEEE, 2022.
- [83] L. An, A.-J. Tu, X. Liu, and R. Akkiraju. Real-time statistical log anomaly detection with continuous aiops learning. In *CLOSER*, pages 223–230, 2022.
- [84] V. Anand, M. Stolet, T. Davidson, I. Beschastnikh, T. Munzner, and J. Mace. Aggregate-driven trace visualizations for performance debugging. *arXiv preprint arXiv:2010.13681*, 2020.
- [85] V. Anand and J. Wonsil. Tracey-distributed trace comparison and aggregation using nlp techniques.
- [86] M. V. d. A. ANDRADE. Migração de microserviços containerizados em tempo de execução. Master’s thesis, Universidade Federal de Pernambuco, 2018.
- [87] T. Angerstein, C. Heger, A. v. Hoorn, D. Okanović, H. Schulz, S. Siegl, and A. Wert. Continuity-automatisiertes performance-testen in der kontinuierlichen softwareentwicklung: Abschlussbericht. 2020.
- [88] V. Arya, K. Shanmugam, P. Aggarwal, Q. Wang, P. Mohapatra, and S. Nagar. Evaluation of causal inference techniques for aiops. In *8th ACM IKDD CODS and 26th COMAD*, pages 188–192. 2021.
- [89] asc lab. Asclab micronaut poc - lab insurance sales portal. <https://github.com/asc-lab/micronaut-microservices-poc/>.
- [90] S. Ashok, P. B. Godfrey, and R. Mittal. Leveraging service meshes as a new network layer. In *Proceedings of the Twentieth ACM Workshop on Hot Topics in Networks*, pages 229–236, 2021.
- [91] W. K. Assunção, J. Krüger, and W. D. Mendonça. Variability management meets microservices: six challenges of re-engineering microservice-based webshops. In *Proceedings of the 24th ACM Conference on Systems and Software Product Line: Volume A-Volume A*, pages 1–6, 2020.
- [92] S. Athlur, N. Sondhi, S. Batra, S. Kalambur, and D. Sitaram. Cache characterization of workloads in a microservice environment. In *2019 IEEE International Conference on Cloud Computing in Emerging Markets (CCEM)*, pages 45–50. IEEE, 2019.
- [93] A. Avritzer, V. Ferme, A. Janes, B. Russo, H. Schulz, and A. van Hoorn. A quantitative approach for the assessment of microservice architecture deployment alternatives by automated performance testing. In *European Conference on Software Architecture*, pages 159–174. Springer, 2018.
- [94] A. Aznavouridis, K. Tsakos, and E. G. Petrakis. Micro-service placement policies for cost optimization in kubernetes. In *International Conference on Advanced Information Networking and Applications*, pages 409–420. Springer, 2022.
- [95] A. F. Baarzi and G. Kesidis. Showar: Right-sizing and efficient scheduling of microservices. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 427–441, 2021.
- [96] A. Baluta, J. Mukherjee, and M. Litoiu. Machine learning based interference modelling in cloud-native applications. In *Proceedings of the 2022 ACM/SPEC on International Conference on Performance Engineering*, pages 125–132, 2022.
- [97] L. Bao, C. Wu, X. Bu, N. Ren, and M. Shen. Performance modeling and workflow scheduling of microservice-based applications in clouds. *IEEE Transactions on Parallel and Distributed Systems*, 30(9):2114–2129, 2019.
- [98] G. C. Baptista et al. *Failure injection in microservice applications*. PhD thesis, Universidade de Coimbra, 2021.
- [99] Barber0. Deathstarbench social network grpc. Retrieved August 2022 from https://github.com/Barber0/DeathStarBench/tree/sn_grpc.
- [100] L. Baresi, D. Y. X. Hu, G. Quattrocchi, and L. Terracciano. Neptune: Network-and gpu-aware management of serverless functions at the edge. *arXiv preprint arXiv:2205.04320*, 2022.
- [101] J. Berg, F. Ruffy, K. Nguyen, N. Yang, T. Kim, A. Sivaraman, R. Ne-travali, and S. Narayana. Snicket: Query-driven distributed tracing. In *Proceedings of the Twentieth ACM Workshop on Hot Topics in Networks*, pages 206–212, 2021.
- [102] I. Beschastnikh, P. Liu, A. Xing, P. Wang, Y. Brun, and M. D. Ernst. Visualizing distributed system executions. *ACM Transactions on Software Engineering and Methodology (TOSEM)*, 29(2):1–38, 2020.
- [103] V. M. Bhasi, J. R. Gunasekaran, P. Thinakaran, C. S. Mishra, M. T. Kandemir, and C. Das. Kraken: Adaptive container provisioning for deploying dynamic dags in serverless platforms. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 153–167, 2021.
- [104] N. Bjørndal, A. Bucchiarone, M. Mazzara, N. Dragoni, S. Dustdar, F. B. Kessler, and T. Wien. Migration from monolith to microservices: Benchmarking a case study. 2020.
- [105] M. Bogo, J. Soldani, D. Neri, and A. Brogi. Fine-grained management of cloud-native applications, based on toasca. *Internet Technology Letters*, 3(5):e212, 2020.
- [106] A. Brogi, A. Corradini, and J. Soldani. Estimating costs of multi-component enterprise applications. *Formal Aspects of Computing*, 31(4):421–451, 2019.
- [107] R. Brondolin and M. D. Santambrogio. A black-box monitoring approach to measure microservices runtime performance. *ACM Transactions on Architecture and Code Optimization (TACO)*, 17(4):1–26, 2020.
- [108] R. Brondolin and M. D. Santambrogio. Presto: a latency-aware power-capping orchestrator for cloud-native microservices. In *2020 IEEE International Conference on Autonomic Computing and Self-Organizing Systems (ACSOS)*, pages 11–20. IEEE, 2020.
- [109] V. Bushong, D. Das, A. Al Maruf, and T. Cerny. Using static analysis to address microservice architecture reconstruction. In *2021 36th IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 1199–1201. IEEE, 2021.
- [110] V. Bushong, D. Das, and T. Cerny. Reconstructing the holistic architecture of microservice systems using static analysis. In *CLOSER*, pages 149–157, 2022.
- [111] A. Byrne, S. Nadgowda, and A. K. Coskun. Ace: Just-in-time serverless software component discovery through approximate concrete execution. In *Proceedings of the 2020 Sixth International Workshop on Serverless Computing*, pages 37–42, 2020.
- [112] S. Caculo, K. Lahiri, and S. Kalambur. Characterizing the scale-up performance of microservices using teastore. In *2020 IEEE International Symposium on Workload Characterization (IISWC)*, pages 48–59. IEEE, 2020.
- [113] M. Camilli, A. Guerriero, A. Janes, B. Russo, and S. Russo. Microservices integrated performance and reliability testing. In *Proceedings of the 3rd ACM/IEEE International Conference on Automation of Software Test*, pages 29–39, 2022.
- [114] M. Camilli, A. Janes, and B. Russo. Automated test-based learning and verification of performance models for microservices systems. *Journal of Systems and Software*, 187:111225, 2022.

- [115] M. Camilli and B. Russo. Modeling performance of microservices systems with growth theory. *Empirical Software Engineering*, 27(2):1–44, 2022.
- [116] L. Cao and P. Sharma. Co-locating containerized workload using service mesh telemetry. In *Proceedings of the 17th International Conference on emerging Networking EXperiments and Technologies*, pages 168–174, 2021.
- [117] W. Cao, Z. Cao, and X. Zhang. Research on microservice anomaly detection technology based on conditional random field. In *Journal of Physics: Conference Series*, volume 1213, page 042016. IOP Publishing, 2019.
- [118] V. Cardellini, E. Casalicchio, S. Iannucci, M. Lucantonio, S. Mittal, D. Panigrahi, and A. Silvi. An intrusion response system utilizing deep q-networks and system partitions. *arXiv preprint arXiv:2202.08182*, 2022.
- [119] A. Carrusca, M. C. Gomes, and J. Leitão. Microservices management on cloud/edge environments. In *International Conference on Service-Oriented Computing*, pages 95–108. Springer, 2019.
- [120] C. Cassé, P. Berthou, P. Owezarski, and S. Josset. Using distributed tracing to identify inefficient resources composition in cloud applications. In *2021 IEEE 10th International Conference on Cloud Networking (CloudNet)*, pages 40–47. IEEE, 2021.
- [121] N. Chen, H. Tu, X. Duan, L. Hu, and C. Guo. Semisupervised anomaly detection of multivariate time series based on a variational autoencoder. *Applied Intelligence*, pages 1–25, 2022.
- [122] X. Chen, H. Irshad, Y. Chen, A. Gehani, and V. Yegneswaran. {CLARION}: Sound and clear provenance tracking for microservice deployments. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 3989–4006, 2021.
- [123] B. Choi, J. Park, C. Lee, and D. Han. phpa: A proactive autoscaling framework for microservice chain. In *5th Asia-Pacific Workshop on Networking (APNet 2021)*, pages 65–71, 2021.
- [124] N. Chondamrongkul, J. Sun, and I. Warren. Formal security analysis for software architecture design: An expressive framework to emerging architectural styles. *Science of Computer Programming*, 206:102631, 2021.
- [125] S. Choochotkaew, T. Chiba, S. Trent, and M. Amaral. Run wild: Resource management system with generalized modeling for microservices on cloud. In *2021 IEEE 14th International Conference on Cloud Computing (CLOUD)*, pages 609–618. IEEE, 2021.
- [126] K.-H. Chow, U. Deshpande, S. Seshadri, and L. Liu. Deeprest: deep resource estimation for interactive microservices. In *Proceedings of the Seventeenth European Conference on Computer Systems*, pages 181–198, 2022.
- [127] V. Cortellessa, D. Di Pompeo, R. Eramo, and M. Tucci. A model-driven approach for continuous performance engineering in microservice-based systems. *Journal of Systems and Software*, 183:111084, 2022.
- [128] V. Cortellessa and L. Traini. Detecting latency degradation patterns in service-based systems. In *Proceedings of the ACM/SPEC International Conference on Performance Engineering*, pages 161–172, 2020.
- [129] C. Courageux-Sudan, A.-C. Orgerie, and M. Quinson. Automated performance prediction of microservice applications using simulation. In *2021 29th International Symposium on Modeling, Analysis, and Simulation of Computer and Telecommunication Systems (MASCOTS)*, pages 1–8. IEEE, 2021.
- [130] C. Cui, G. Wu, W. Chen, J. Zhu, and J. Wei. Feedback-based, automated failure testing of microservice-based applications. *arXiv preprint arXiv:1908.06466*, 2019.
- [131] J. Curtis. On language-agnostic abstract-syntax trees: student research abstract. In *Proceedings of the 37th ACM/SIGAPP Symposium on Applied Computing*, pages 1619–1625, 2022.
- [132] Q. N. Dao and T. Hey. *Coreference Resolution for Software Architecture Documentation*. PhD thesis, Informatics Institute, 2022.
- [133] N. Daw, U. Bellur, and P. Kulkarni. Speedo: Fast dispatch and orchestration of serverless workflows. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 585–599, 2021.
- [134] E. De Angelis, A. Pellegrini, and M. Proietti. Automatic extraction of behavioral features for test program similarity analysis. In *2021 IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW)*, pages 129–136. IEEE, 2021.
- [135] A. De Iasio and E. Zimeo. A framework for microservices synchronization. *Software: Practice and Experience*, 51(1):25–45, 2021.
- [136] U. Desai, S. Bandyopadhyay, and S. Tamilselvam. Graph neural network to dilute outliers for refactoring monolith application. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 72–80, 2021.
- [137] D. Di Pompeo, M. Tucci, A. Celi, and R. Eramo. A microservice reference case study for design-runtime interaction in mde. In *STAF (Co-Located Events)*, pages 23–32, 2019.
- [138] M. Dinga. An empirical evaluation of the energy and performance overhead of monitoring tools on docker-based systems. Retrieved August 2022 from <https://github.com/MadalinaDinga/thesis-2022-monitoring-tools-integration-with-train-ticket-replication-package/>, 2022.
- [139] P. Dogga, K. Narasimhan, A. Sivaraman, S. Saini, G. Varghese, and R. Netravali. Revelio: ML-generated debugging queries for finding root causes in distributed systems. *Proceedings of Machine Learning and Systems*, 4:601–622, 2022.
- [140] dot-net architecture. eshoponcontainers. <https://github.com/dotnet-architecture/eShopOnContainers>.
- [141] T. F. Düllmann, A. van Hoorn, V. Yussupov, P. Jakovits, and M. Adhikari. Ctt: Load test automation for toasca-based cloud applications. In *Companion of the 2022 ACM/SPEC International Conference on Performance Engineering*, pages 89–96, 2022.
- [142] EdwinVW. Pitstop: Garage management system. <https://github.com/EdwinVW/pitstop/>.
- [143] S. Eismann, C.-P. Bezemer, W. Shang, D. Okanović, and A. van Hoorn. Microservices: A performance tester’s dream or nightmare? In *Proceedings of the ACM/SPEC International Conference on Performance Engineering*, pages 138–149, 2020.
- [144] A. El Malki and U. Zdun. Evaluation of api request bundling and its impact on performance of microservice architectures. In *2021 IEEE International Conference on Services Computing (SCC)*, pages 419–424. IEEE, 2021.
- [145] S. Elflein, I. S. Kounev, J. Grohmann, and S. Eismann. Automatic identification of parametric dependencies for performance models.
- [146] M. Elsaadawy, M. Younis, J. Wang, X. Hou, and B. Kemme. Dynamic application call graph formation and service identification in cloud data centers. *IEEE Transactions on Network and Service Management*, 2022.
- [147] D. Elsner, D. Bertagnolli, A. Pretschner, and R. Klaus. Challenges in regression test selection for end-to-end testing of microservice-based software systems. In *International Conference on Automation of Software Test (AST’22)*, 2022.
- [148] A. Erradi, W. Iqbal, A. Mahmood, and A. Bouguettaya. Web application resource requirements estimation based on the workload latent features. *IEEE Transactions on Services Computing*, 2019.
- [149] A. M. Espinoza, R. Wood, S. Forrest, and M. Tiwari. Back to the future: N-versioning of microservices. In *2022 52nd Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*, pages 415–427. IEEE, 2022.
- [150] A. Farley. Continuous performance testing of faas and microservices based on toasca topology and orchestration specifications. B.S. thesis, 2020.
- [151] J. Faro. Deathstarbench metastability failure fork. Retrieved August 2022 from <https://github.com/jfaro/DeathStarBench>, 2022.

- [152] J. Faro. Metastable failure vulnerability testing for deathstarbench's social network micro-service benchmark. Retrieved August 2022 from <https://github.com/jfaro/metastable-failures>, 2022.
- [153] K. Fedyuk. *Sheik: Dynamic location and binding of microservices for cloud/edge settings*. PhD thesis, NOVA University of Lisbon, 2020.
- [154] J. Flora, P. Gonçalves, M. Teixeira, and N. Antunes. My services got old! can kubernetes handle the aging of microservices? In *2021 IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW)*, pages 40–47. IEEE, 2021.
- [155] J. Forough, M. Bhuyan, and E. Elmroth. Detection of vsi-ddos attacks on the edge: A sequential modeling approach. In *The 16th International Conference on Availability, Reliability and Security*, pages 1–10, 2021.
- [156] K. Fu, W. Zhang, Q. Chen, D. Zeng, and M. Guo. Adaptive resource efficient microservice deployment in cloud-edge continuum. *IEEE Transactions on Parallel and Distributed Systems*, 33(8):1825–1840, 2021.
- [157] A. Fuerst, A. Ali-Eldin, P. Shenoy, and P. Sharma. Cloud-scale vm-deflation for running interactive applications on transient servers. In *Proceedings of the 29th International Symposium on High-Performance Parallel and Distributed Computing*, pages 53–64, 2020.
- [158] N. Fukuda, C. Wu, S. Horiuchi, and K. Tayama. Fault report generation for ict systems by jointly learning time-series and text data. In *NOMS 2022-2022 IEEE/IFIP Network Operations and Management Symposium*, pages 1–9. IEEE, 2022.
- [159] S. Galantino, M. Iorio, F. Risso, and A. Manzalini. Raygo: Reserve as you go. In *2021 IEEE Intl Conf on Dependable, Autonomic and Secure Computing, Intl Conf on Pervasive Intelligence and Computing, Intl Conf on Cloud and Big Data Computing, Intl Conf on Cyber Science and Technology Congress (DASC/PiCom/CBDCom/CyberSciTech)*, pages 269–275. IEEE, 2021.
- [160] Y. Gan, M. Liang, D. L. Sundar Dev, and C. Delimitrou. Sage: Practical & scalable ml-driven performance debugging in microservices.
- [161] Y. Gan, Y. Zhang, D. Cheng, A. Shetty, P. Rathi, N. Katarki, A. Bruno, J. Hu, B. Ritchken, B. Jackson, et al. An open-source benchmark suite for microservices and their hardware-software implications for cloud & edge systems. In *Proceedings of the Twenty-Fourth International Conference on Architectural Support for Programming Languages and Operating Systems*, pages 3–18, 2019.
- [162] Y. Gan, Y. Zhang, K. Hu, D. Cheng, Y. He, M. Pancholi, and C. Delimitrou. Seer: Leveraging big data to navigate the complexity of performance debugging in cloud microservices. In *Proceedings of the twenty-fourth international conference on architectural support for programming languages and operating systems*, pages 19–33, 2019.
- [163] D. Garg, N. C. Narendra, and S. Tesfatsion. Heuristic and reinforcement learning algorithms for dynamic service placement on mobile edge cloud. *arXiv preprint arXiv:2111.00240*, 2021.
- [164] P. Genfer and U. Zdun. Identifying domain-based cyclic dependencies in microservice apis using source code detectors. In *European Conference on Software Architecture*, pages 207–222. Springer, 2021.
- [165] P. Genfer and U. Zdun. Avoiding excessive data exposure through microservice apis. In *European Conference on Software Architecture*, pages 3–18. Springer, 2022.
- [166] S. Gholami, A. Goli, C.-P. Bezemer, and H. Khazaei. A framework for satisfying the performance requirements of containerized software systems through multi-versioning. In *Proceedings of the ACM/SPEC International Conference on Performance Engineering*, pages 150–160, 2020.
- [167] A. Gluck. Introducing domain-oriented microservice architecture. (July 2020). Retrieved October 2020 from <https://eng.uber.com/microservice-architecture/>, 2020.
- [168] A. Goli, N. Mahmoudi, H. Khazaei, and O. Ardakanian. A holistic machine learning-based autoscaling approach for microservice applications. In *CLOSER*, pages 190–198, 2021.
- [169] GoogleCloudPlatform. Online boutique, fka hipstershop. <https://github.com/GoogleCloudPlatform/microservices-demo>.
- [170] M. Grambow, T. Pfandzelter, L. Burchard, C. Schubert, M. Zhao, and D. Bermbach. Befaa: An application-centric benchmarking framework for faas platforms. In *2021 IEEE International Conference on Cloud Engineering (IC2E)*, pages 1–8. IEEE, 2021.
- [171] M. Grambow, E. Wittern, and D. Bermbach. Benchmarking the performance of microservice applications. *ACM SIGAPP Applied Computing Review*, 20(3):20–34, 2020.
- [172] D. Granata, M. Rak, and G. Salzillo. Metasend: A security enabled development life cycle meta-model. In *Proceedings of the 17th International Conference on Availability, Reliability and Security*, pages 1–10, 2022.
- [173] G. Granchelli, M. Cardarelli, P. Di Francesco, I. Malavolta, L. Iovino, and A. Di Salle. Microart: A software architecture recovery tool for maintaining microservice-based systems. In *2017 IEEE International Conference on Software Architecture Workshops (ICSAW)*, pages 298–302. IEEE, 2017.
- [174] G. Granchelli, M. Cardarelli, P. Di Francesco, I. Malavolta, L. Iovino, and A. Di Salle. Towards recovering the software architecture of microservice-based systems. In *2017 IEEE International Conference on Software Architecture Workshops (ICSAW)*, pages 46–53. IEEE, 2017.
- [175] J. Grohmann, S. Eismann, S. Elflein, J. v. Kistowski, S. Kounev, and M. Mazkati. Detecting parametric dependencies for performance models using feature selection techniques. In *2019 IEEE 27th International Symposium on Modeling, Analysis, and Simulation of Computer and Telecommunication Systems (MASCOTS)*, pages 309–322. IEEE, 2019.
- [176] J. Grohmann, P. K. Nicholson, J. O. Iglesias, S. Kounev, and D. Lugones. Monitorless: Predicting performance degradation in cloud applications with machine learning. In *Proceedings of the 20th international middleware conference*, pages 149–162, 2019.
- [177] J. Grohmann, M. Straesser, A. Chalbani, S. Eismann, Y. Arian, N. Herbst, N. Peretz, and S. Kounev. Suanming: Explainable prediction of performance degradations in microservice applications. In *Proceedings of the ACM/SPEC International Conference on Performance Engineering*, pages 165–176, 2021.
- [178] Z. Guan, J. Lin, and P. Chen. On anomaly detection and root cause analysis of microservice systems. In *International Conference on Service-Oriented Computing*, pages 465–469. Springer, 2018.
- [179] C. Guerrero, I. Lera, and C. Juiz. Genetic algorithm for multi-objective optimization of container allocation in cloud architecture. *Journal of Grid Computing*, 16(1):113–135, 2018.
- [180] C. Guerrero, I. Lera, and C. Juiz. Evaluation and efficiency comparison of evolutionary algorithms for service placement optimization in fog architectures. *Future Generation Computer Systems*, 97:131–144, 2019.
- [181] C. Guerrero, I. Lera, and C. Juiz. A lightweight decentralized service placement policy for performance optimization in fog computing. *Journal of Ambient Intelligence and Humanized Computing*, 10(6):2435–2452, 2019.
- [182] A. Guerriero and L. Giamattei. Microsertest: a flexible tool for automated testing of microservices applications.
- [183] D. Guhathakurta, P. Aggarwal, S. Nagar, R. Arora, and B. Zhou. Utilizing persistence for post facto suppression of invalid anomalies using system logs. 2022.
- [184] A. Hakamian. Engineering resilience: Predicting the change impact on performance and availability of reconfigurable systems. In *2020 IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW)*, pages 143–146. IEEE, 2020.
- [185] W. Hasselbring. Benchmarking as empirical standard in software engineering research. In *Evaluation and Assessment in Software Engineering*, pages 365–372. 2021.
- [186] Z. He, P. Chen, X. Li, Y. Wang, G. Yu, C. Chen, X. Li, and Z. Zheng. A spatiotemporal deep learning approach for unsupervised anomaly

- detection in cloud systems. *IEEE Transactions on Neural Networks and Learning Systems*, 2020.
- [187] A. Horn, H. M. Fard, and F. Wolf. Multi-objective hybrid autoscaling of microservices in kubernetes clusters. In *European Conference on Parallel Processing*, pages 233–250. Springer, 2022.
 - [188] M. R. Hossen and M. A. Islam. Practical efficient microservice autoscaling.
 - [189] M. R. Hossen and M. A. Islam. Towards efficient microservices management through opportunistic resource reduction.
 - [190] M. R. Hossen and M. A. Islam. A lightweight workload-aware microservices autoscaling with qos assurance. *arXiv preprint arXiv:2202.00057*, 2022.
 - [191] X. Hou, C. Li, J. Liu, L. Zhang, Y. Hu, and M. Guo. Ant-man: towards agile power management in the microservice era. In *2020 SC20: International Conference for High Performance Computing, Networking, Storage and Analysis (SC)*, pages 1098–1111. IEEE Computer Society, 2020.
 - [192] X. Hou, J. Liu, C. Li, and M. Guo. Unleashing the scalability potential of power-constrained data center in the microservice era. In *Proceedings of the 48th International Conference on Parallel Processing*, pages 1–10, 2019.
 - [193] L. Huang and T. Zhu. tprof: Performance profiling via structural aggregation and automated analysis of distributed systems traces. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 76–91, 2021.
 - [194] L. Iffländer, N. Rawtani, L. Beierlieb, N. Fella, K.-D. Lange, and S. Kounev. Attack-aware security function chain reordering. *arXiv preprint arXiv:2005.08364*, 2020.
 - [195] T. Inagaki, Y. Ueda, T. Nakaike, and M. Ohara. Profile-based detection of layered bottlenecks. In *Proceedings of the 2019 ACM/SPEC International Conference on Performance Engineering*, pages 197–208, 2019.
 - [196] T. Inagaki, Y. Ueda, M. Ohara, S. Choochotkaew, M. Amaral, S. Trent, T. Chiba, and Q. Zhang. Detecting layered bottlenecks in microservices. In *2022 IEEE 15th International Conference on Cloud Computing (CLOUD)*, pages 385–396. IEEE, 2022.
 - [197] M. R. Islam, A. Al Maruf, and T. Cerny. Code smell prioritization with business process mining and static code analysis: A case study. *Electronics*, 11(12):1880, 2022.
 - [198] A. A. T. Isstaif. Self-managed services using mirageos unikernels. In *Proceedings of the 21st International Middleware Conference Doctoral Symposium*, pages 35–37, 2020.
 - [199] istio. Bookinfo application. <https://istio.io/latest/docs/examples/bookinfo/>.
 - [200] J. Ivers, C. Seifried, and I. Ozkaya. Untangling the knot: Enabling architecture evolution with search-based refactoring. In *2022 IEEE 19th International Conference on Software Architecture (ICSA)*, pages 101–111. IEEE, 2022.
 - [201] S. Jacob, Y. Qiao, and B. Lee. Detecting cyber security attacks against a microservices application using distributed tracing. In *ICISSP*, pages 588–595, 2021.
 - [202] S. Jacob, Y. Qiao, Y. Ye, and B. Lee. Anomalous distributed traffic: Detecting cyber security attacks amongst microservices using graph convolutional networks. *Computers & Security*, 118:102728, 2022.
 - [203] S. Ji, W. Wu, and Y. Pu. Multi-indicators prediction in microservice using granger causality test and attention lstm. In *2020 IEEE World Congress on Services (SERVICES)*, pages 77–82. IEEE, 2020.
 - [204] T. Jia, Y. Li, C. Zhang, W. Xia, J. Jiang, and Y. Liu. Machine deserves better logging: a log enhancement approach for automatic fault diagnosis. In *2018 IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW)*, pages 106–111. IEEE, 2018.
 - [205] Z. Jia and E. Witchel. Nightcore: Efficient and scalable serverless computing for latency-sensitive, interactive microservices extended abstract.
 - [206] Z. Jin, Y. Zhu, J. Zhu, D. Yu, C. Li, R. Chen, I. E. Akkus, and Y. Xu. Lessons learned from migrating complex stateful applications onto serverless platforms. In *Proceedings of the 12th ACM SIGOPS Asia-Pacific Workshop on Systems*, pages 89–96, 2021.
 - [207] D. Johansson. Model-driven development for microservices: A domain-specific modeling language for kubernetes, 2022.
 - [208] C. T. Joseph and K. Chandrasekaran. Intma: Dynamic interaction-aware resource allocation for containerized microservices in cloud environments. *Journal of Systems Architecture*, 111:101785, 2020.
 - [209] H. Kadiyala, A. Misail, and J. Rubin. Kuber: cost-efficient microservice deployment planner. In *2022 IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER)*, pages 252–262. IEEE, 2022.
 - [210] M. Kang and J. Kim. Comparison of service description and composition for complex 3-tier cloud-based services. *Proceedings of the Asia-Pacific Advanced Network*, 44:37–40, 2017.
 - [211] M. Kang, J.-S. Shin, and J. Kim. Protected coordination of service mesh for container-based 3-tier service traffic. In *2019 International Conference on Information Networking (ICOIN)*, pages 427–429. IEEE, 2019.
 - [212] R. R. Karn, R. Das, D. R. Pant, J. Heikkonen, and R. Kanth. Automated testing and resilience of microservice’s network-link using istio service mesh. In *2022 31st Conference of Open Innovations Association (FRUCT)*, pages 79–88. IEEE, 2022.
 - [213] J. Keim, S. Schulz, D. Fuchß, C. Kocher, J. Speit, and A. Koziol. Trace link recovery for software architecture documentation. In *European Conference on Software Architecture*, pages 101–116. Springer, 2021.
 - [214] A. Khrabrov, M. Pirvu, V. Sundaresan, and E. De Lara. {JITServer}: Disaggregated caching {JIT} compiler for the {JVM} in the cloud. In *2022 USENIX Annual Technical Conference (USENIX ATC 22)*, pages 869–884, 2022.
 - [215] P. Kokić. *Unaprijedjenje sigurnosti pomoću mikroservisa*. PhD thesis, University of Zagreb. Faculty of Organization and Informatics. Department of ..., 2018.
 - [216] J. Kosińska and K. Zieliński. Autonomic management framework for cloud-native applications. *Journal of Grid Computing*, 18(4):779–796, 2020.
 - [217] P. Kraft, Q. Li, K. Kaffes, A. Skiadopoulos, D. Kumar, D. Cho, J. Li, R. Redmond, N. Weckwerth, B. Xia, et al. Apiary: A dbms-backed transactional function-as-a-service framework. *arXiv preprint arXiv:2208.13068*, 2022.
 - [218] N. Kratzke. About the complexity to transfer cloud applications at runtime and how container platforms can contribute? In *International Conference on Cloud Computing and Services Science*, pages 19–45. Springer, 2017.
 - [219] N. Kratzke. Smuggling multi-cloud support into cloud-native applications using elastic container platforms. In *CLOSER*, pages 29–42, 2017.
 - [220] N. Kratzke. About an immune system understanding for cloud-native applications. *CLOUD COMPUTING 2018*, page 41, 2018.
 - [221] N. Kratzke. About being the tortoise or the hare?-a position paper on making cloud applications too fast and furious for attackers. *arXiv preprint arXiv:1802.03565*, 2018.
 - [222] N. Kratzke and P.-C. Quint. Project cloudtransit.
 - [223] R. Laigner, Y. Zhou, and M. A. V. Salles. A distributed database system for event-based microservices. In *Proceedings of the 15th ACM International Conference on Distributed and Event-based Systems*, pages 25–30, 2021.
 - [224] P. Las-Casas, G. Papakerashvili, V. Anand, and J. Mace. Sifter: Scalable sampling for distributed traces, without feature engineering. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 312–324, 2019.
 - [225] N. Lazarev, N. Adit, S. Xiang, Z. Zhang, and C. Delimitrou. Dagger: Towards efficient rpcs in cloud microservices with near-memory

- reconfigurable nics. *IEEE Computer Architecture Letters*, 19(2):134–138, 2020.
- [226] A. Lesniak, R. Laigner, and Y. Zhou. Enforcing consistency in microservice architectures through event-based constraints. In *Proceedings of the 15th ACM International Conference on Distributed and Event-based Systems*, pages 180–183, 2021.
- [227] J. Levin and T. A. Benson. Viperprobe: Rethinking microservice observability with ebpf. In *2020 IEEE 9th International Conference on Cloud Networking (CloudNet)*, pages 1–8. IEEE, 2020.
- [228] Q. Li, B. Li, P. Mercati, R. Illikkal, C. Tai, M. Kishinevsky, and C. Kozyrakis. Rambo: Resource allocation for microservices using bayesian optimization. *IEEE Computer Architecture Letters*, 20(1):46–49, 2021.
- [229] R. Li, M. Du, H. Chang, S. Mukherjee, and E. Eide. Deepstitch: Deep learning for cross-layer stitching in microservices. In *Proceedings of the 2020 6th International Workshop on Container Technologies and Container Clouds*, pages 25–30, 2020.
- [230] R. Li, M. Du, Z. Wang, H. Chang, S. Mukherjee, and E. Eide. Longtale: Toward automatic performance anomaly explanation in microservices. In *Proceedings of the 2022 ACM/SPEC on International Conference on Performance Engineering*, pages 5–16, 2022.
- [231] X. Li, Y. Chen, and Z. Lin. Towards automated inter-service authorization for microservice applications. In *Proceedings of the ACM SIGCOMM 2019 Conference Posters and Demos*, pages 3–5, 2019.
- [232] X. Li, Y. Chen, Z. Lin, X. Wang, and J. H. Chen. Automatic policy generation for {Inter-Service} access control of microservices. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 3971–3988, 2021.
- [233] Y. Li, T. Li, P. Shen, L. Hao, W. Liu, S. Wang, Y. Song, and L. Bao. Sim-drs: a similarity-based dynamic resource scheduling algorithm for microservice-based web systems. *PeerJ Computer Science*, 7:e824, 2021.
- [234] Z. Li, J. Chen, R. Jiao, N. Zhao, Z. Wang, S. Zhang, Y. Wu, L. Jiang, L. Yan, Z. Wang, et al. Practical root cause localization for microservice systems via trace analysis. In *2021 IEEE/ACM 29th International Symposium on Quality of Service (IWQOS)*, pages 1–10. IEEE, 2021.
- [235] Z. Li, N. Zhao, M. Li, X. Lu, L. Wang, D. Chang, X. Nie, L. Cao, W. Zhang, K. Sui, et al. Actionable and interpretable fault localization for recurring failures in online service systems. *arXiv preprint arXiv:2207.09021*, 2022.
- [236] H. Lim, Y. Kim, and K. Sun. Service management in virtual machine and container mixed environment using service mesh. In *2021 International Conference on Information Networking (ICOIN)*, pages 528–530. IEEE, 2021.
- [237] F. Liu, Y. Wang, Z. Li, R. Ren, H. Guan, X. Yu, X. Chen, and G. Xie. Microcbr: Case-based reasoning on spatio-temporal fault knowledge graph for microservices troubleshooting. In *International Conference on Case-Based Reasoning*, pages 224–239. Springer, 2022.
- [238] J. Loff. Deathstarbench inconsistency. Retrieved August 2022 from <https://github.com/jfloff/antipode-deathstarbench>, 2022.
- [239] Z. Long, G. Wu, X. Chen, C. Cui, W. Chen, and J. Wei. Fitness-guided resilience testing of microservice-based applications. In *2020 IEEE International Conference on Web Services (ICWS)*, pages 151–158. IEEE, 2020.
- [240] N. D. S. Loureiro. *Compositional analysis of vulnerabilities in microservice-based applications*. PhD thesis, 2021.
- [241] S. Luo, H. Xu, C. Lu, K. Ye, G. Xu, L. Zhang, Y. Ding, J. He, and C. Xu. Characterizing microservice dependency and performance: Alibaba trace analysis. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 412–426, 2021.
- [242] S. Luo, H. Xu, C. Lu, K. Ye, G. Xu, L. Zhang, J. He, and C. Xu. An in-depth study of microservice call graph and runtime performance. *IEEE Transactions on Parallel and Distributed Systems*, 33(12):3901–3914, 2022.
- [243] W. Lv, Q. Wang, P. Yang, Y. Ding, B. Yi, Z. Wang, and C. Lin. Microservice deployment in edge computing based on deep q learning. *IEEE Transactions on Parallel and Distributed Systems*, 2022.
- [244] J. Mace and R. Fonseca. Universal context propagation for distributed system instrumentation. In *Proceedings of the Thirteenth EuroSys Conference*, pages 1–18, 2018.
- [245] N. Marie-Magdelaine and T. Ahmed. Proactive autoscaling for cloud-native applications using machine learning. In *GLOBECOM 2020-2020 IEEE Global Communications Conference*, pages 1–7. IEEE, 2020.
- [246] D. R. Mathews, M. Verma, J. Lakshmi, and P. Aggarwal. Towards more effective and explainable fault management using cross-layer service topology. In *2022 IEEE 15th International Conference on Cloud Computing (CLOUD)*, pages 94–96. IEEE, 2022.
- [247] A. P. Matthews. *Microservice Pattern Identification from Recovered Architectures of Orchestrated Systems*. University of California, Irvine, 2021.
- [248] M. Mazkatli, D. Monschein, J. Grohmann, and A. Koziolk. Incremental calibration of architectural performance models with parametric dependencies. In *2020 IEEE International Conference on Software Architecture (ICSA)*, pages 23–34. IEEE, 2020.
- [249] D. Mealha, N. Preguiça, M. C. Gomes, and J. Leitão. Data replication on the cloud/edge. In *Proceedings of the 6th Workshop on Principles and Practice of Consistency for Distributed Data*, pages 1–7, 2019.
- [250] W. D. Mendonça, W. K. Assunção, L. V. Estanislau, S. R. Vergilio, and A. Garcia. Towards a microservices-based product line with multi-objective evolutionary algorithms. In *2020 IEEE Congress on Evolutionary Computation (CEC)*, pages 1–8. IEEE, 2020.
- [251] microservices demo. Retrieved August 2022 from <https://github.com/microservices-demo/microservices-demo>, 2016.
- [252] A. Mirhosseini, S. Elnikety, and T. F. Wenisch. Parslo: A gradient descent-based approach for near-optimal partial slo allotment in microservices. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 442–457, 2021.
- [253] A. Mirhosseini, A. Sriraman, and T. F. Wenisch. Enhancing server efficiency in the face of killer microseconds. In *2019 IEEE International Symposium on High Performance Computer Architecture (HPCA)*, pages 185–198. IEEE, 2019.
- [254] H. Mohamed and O. El-Gayar. End-to-end latency prediction of microservices workflow on kubernetes: A comparative evaluation of machine learning models and resource metrics. In *Proceedings of the 54th Hawaii International Conference on System Sciences*, page 1717, 2021.
- [255] J. Mukherjee, A. Baluta, M. Litoiu, and D. Krishnamurthy. Rad: Detecting performance anomalies in cloud-based web services. In *2020 IEEE 13th International Conference on Cloud Computing (CLOUD)*, pages 493–501. IEEE, 2020.
- [256] G. Muntoni, J. Soldani, and A. Brogi. Mining the architecture of microservice-based applications from their kubernetes deployment. In *European Conference on Service-Oriented and Cloud Computing*, pages 103–115. Springer, 2020.
- [257] A. Nadeem and M. Z. Malik. A case for microservices orchestration using workflow engines. *arXiv preprint arXiv:2204.07210*, 2022.
- [258] S. Nagar, S. Samanta, P. Mohapatra, and D. Kar. Building golden signal based signatures for log anomaly detection. In *2022 IEEE 15th International Conference on Cloud Computing (CLOUD)*, pages 203–208. IEEE, 2022.
- [259] J. Nam, S. Lee, H. Seo, P. Porras, V. Yegneswaran, and S. Shin. {BASTION}: A security enforcement network stack for container networks. In *2020 {USENIX} Annual Technical Conference ({USENIX} {ATC} 20)*, pages 81–95, 2020.
- [260] A. R. Nasab, M. Shahin, S. A. H. Raviz, P. Liang, A. Mashmool, and V. Lenarduzzi. An empirical study of security practices for microservices systems. *arXiv preprint arXiv:2112.14927*, 2021.

- [261] T. Nebaeus. Optimal scaling configurations for microservice-oriented architectures using genetic algorithms, 2019.
- [262] C. Nguyen, A. Mehta, C. Klein, and E. Elmroth. Why cloud applications are not ready for the edge (yet). In *Proceedings of the 4th ACM/IEEE Symposium on Edge Computing*, pages 250–263, 2019.
- [263] H. X. Nguyen, S. Zhu, and M. Liu. Graph-phpa: Graph-based proactive horizontal pod autoscaling for microservices using lstm-gnn. *arXiv preprint arXiv:2209.02551*, 2022.
- [264] S. Nijboer. Traffic analysis of a service mesh. Master’s thesis, 2021.
- [265] E. T. Nordli, P. H. Nguyen, F. Chauvel, and H. Song. Event-based customization of multi-tenant saas using microservices. In *International Conference on Coordination Languages and Models*, pages 171–180. Springer, 2020.
- [266] E. Ntontos, U. Zdun, K. Plakidas, P. Genfer, S. Geiger, S. Meixner, and W. Hasselbring. Detector-based component model abstraction for microservice-based systems. *Computing*, 103(11):2521–2551, 2021.
- [267] E. Ntontos, U. Zdun, K. Plakidas, S. Meixner, and S. Geiger. Metrics for assessing architecture conformance to microservice architecture patterns and practices. In *International Conference on Service-Oriented Computing*, pages 580–596. Springer, 2020.
- [268] E. Ntontos, U. Zdun, J. Soldani, and A. Brogi. Assessing architecture conformance to coupling-related infrastructure-as-code best practices: Metrics and case studies. In *European Conference on Software Architecture*, pages 101–116. Springer, 2022.
- [269] M. S. Nyfløtt. Optimizing inter-service communication between microservices. Master’s thesis, NTNU, 2017.
- [270] V. Olteanu, H. Eran, D. Dumitrescu, A. Popa, C. Baci, M. Silberstein, G. Nikolaidis, M. Handley, and C. Raiciu. An edge-queued datagram service for all datacenter traffic. In *19th USENIX Symposium on Networked Systems Design and Implementation (NSDI 22)*, pages 761–777, 2022.
- [271] A. Osman, J. Born, and T. Strufe. Mitigating internal, stealthy dos attacks in microservice networks. In *International Symposium on Stabilizing, Safety, and Security of Distributed Systems*, pages 500–504. Springer, 2021.
- [272] overleaf. Overleaf. <https://github.com/overleaf/overleaf>.
- [273] M. Pancholi, A. D. Kellas, V. P. Kemerlis, and S. Sethumadhavan. Timeloops: System call policy learning for containerized microservices. *arXiv preprint arXiv:2204.06131*, 2022.
- [274] J. Park, B. Choi, C. Lee, and D. Han. Graf: a graph neural network based proactive resource allocation framework for slo-oriented microservices. In *Proceedings of the 17th International Conference on emerging Networking EXperiments and Technologies*, pages 154–167, 2021.
- [275] M. A. Patterns. Lakeside mutual. <https://github.com/Microservice-API-Patterns/LakesideMutual>.
- [276] H. Phalachandra, P. Bhatt, I. Agarwal, and R. Bhowmick. Improving task scheduling in microservice environments by considering intra-job dependencies. In *2022 6th International Conference on Intelligent Computing and Control Systems (ICICCS)*, pages 568–573. IEEE, 2022.
- [277] A. Pourhabibi, S. Gupta, H. Kassir, M. Sutherland, Z. Tian, M. P. Drumond, B. Falsafi, and C. Koch. Optimus prime: Accelerating data transformation in servers. In *Proceedings of the Twenty-Fifth International Conference on Architectural Support for Programming Languages and Operating Systems*, pages 1203–1216, 2020.
- [278] A. Pourhabibi, M. Sutherland, A. Daglis, and B. Falsafi. Cerebros: Evading the rpc tax in datacenters. In *MICRO-54: 54th Annual IEEE/ACM International Symposium on Microarchitecture*, pages 407–420, 2021.
- [279] H. Qiu, S. S. Banerjee, S. Jha, Z. T. Kalbarczyk, and R. K. Iyer. {FIRM}: An intelligent fine-grained resource management framework for slo-oriented microservices. In *14th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 20)*, pages 805–825, 2020.
- [280] P.-C. Quint and N. Kratzke. Towards a description of elastic cloud-native applications for transferable multi-cloud-deployments. *Small*, 147:34, 2017.
- [281] P.-C. Quint and N. Kratzke. Towards a lightweight multi-cloud dsl for elastic and transferable cloud-native applications. *arXiv preprint arXiv:1802.03562*, 2018.
- [282] F. Rademacher, S. Sachweh, and A. Zündorf. A modeling method for systematic architecture reconstruction of microservice-based software systems. In *Enterprise, Business-Process and Information Systems Modeling*, pages 311–326. Springer, 2020.
- [283] J. Rahman and P. Lama. Predicting the end-to-end tail latency of containerized microservices in the cloud. In *2019 IEEE International Conference on Cloud Engineering (IC2E)*, pages 200–210. IEEE, 2019.
- [284] M. I. Rahman, S. Panichella, and D. Taibi. A curated dataset of microservices-based systems, extended version. Vol-2520, 2019.
- [285] K. Ramasubramanian, A. Raina, J. Mace, and P. Alvaro. Act now: Aggregate comparison of traces for incident localization. *arXiv preprint arXiv:2205.06933*, 2022.
- [286] V. Rao, V. Singh, K. Goutham, B. U. Kempaiah, R. J. Mampilli, S. Kalambar, and D. Sitaram. Scheduling microservice containers on large core machines through placement and coalescing. In *Workshop on Job Scheduling Strategies for Parallel Processing*, pages 80–100. Springer, 2021.
- [287] K. Ray and A. Banerjee. Horizontal auto-scaling for multi-access edge computing using safe reinforcement learning. *ACM Transactions on Embedded Computing Systems (TECS)*, 20(6):1–33, 2021.
- [288] redis. Redis patterns example. <https://redis.io/docs/reference/patterns/twitter-clone/>.
- [289] J. S. Rellermeyer, M. Amer, R. Smutzer, and K. Rajamani. Container density improvements with dynamic memory extension using nand flash. In *Proceedings of the 9th Asia-Pacific Workshop on Systems*, pages 1–7, 2018.
- [290] F. G. O. Risso and S. Galantino. Enabling job aware scheduling on kubernetes cluster.
- [291] rnp68. Teastore cve-2021-44228. Retrieved September 2022 from <https://github.com/rnp68/TeaStore>, 2021.
- [292] O. R. C. Rodríguez, C. Pahl, N. El Ioini, H. R. Barzegar, et al. Improvement of edge computing workload placement using multi objective particle swarm optimization. In *2021 8th International Conference on Internet of Things: Systems, Management and Security (IOTSMS)*, pages 1–8. IEEE, 2021.
- [293] Y. Rouf, J. Mukherjee, M. Litoiu, J. Wigglesworth, and R. Mateescu. A framework for developing devops operation automation in clouds using components-off-the-shelf. In *Proceedings of the ACM/SPEC International Conference on Performance Engineering*, pages 265–276, 2021.
- [294] W. Runan, G. Casale, and A. Filieri. Service distribution estimation for microservices using markovian arrival processes.
- [295] E. Rwemigabo. *Autonomic application topology redistribution*. PhD thesis, 2018.
- [296] V. Sachidananda and A. Sivaraman. Learned autoscaling for cloud microservices with multi-armed bandits. *arXiv preprint arXiv:2112.14845*, 2021.
- [297] M. Sakai and K. Takahashi. Constructing a service process model based on distributed tracing for conformance checking of microservices. In *NOMS 2022-2022 IEEE/IFIP Network Operations and Management Symposium*, pages 1–6. IEEE, 2022.
- [298] M. Sakai, K. Takahashi, and S. Kondoh. Method of constructing petri net service model using distributed trace data of microservices. In *2021 22nd Asia-Pacific Network Operations and Management Symposium (APNOMS)*, pages 214–217. IEEE, 2021.
- [299] M. R. Saleh Sedghpour, C. Klein, and J. Tordsson. An empirical study of service mesh traffic management policies for microservices. In *Proceedings of the 2022 ACM/SPEC on International Conference on Performance Engineering*, pages 17–27, 2022.
- [300] A. R. SAMPAIO JUNIOR. Runtime adaptation of microservices. 2018.

- [301] M. Santana, A. Sampaio Jr, M. Andrade, and N. S. Rosa. Transparent tracing of microservice-based applications. In *Proceedings of the 34th ACM/SIGAPP Symposium on Applied Computing*, pages 1252–1259, 2019.
- [302] D. Schall, A. Margaritov, D. Ustiugov, A. Sandberg, and B. Grot. Luke-warm serverless functions: characterization and optimization. In *ISCA*, pages 757–770, 2022.
- [303] M. Schiedermeier. A concern-oriented software engineering methodology for micro-service architectures. In *Proceedings of the 23rd ACM/IEEE International Conference on Model Driven Engineering Languages and Systems: Companion Proceedings*, pages 1–5, 2020.
- [304] M. Schiewe. Bridging the gap between source code and high-level concepts in static code analysis: student research abstract. In *Proceedings of the 37th ACM/SIGAPP Symposium on Applied Computing*, pages 1615–1618, 2022.
- [305] M. Schiewe, J. Curtis, V. Bushong, and T. Cerny. Advancing static code analysis with language-agnostic component identification. *IEEE Access*, 10:30743–30761, 2022.
- [306] N. Schmitt, L. Iffländer, A. Bauer, and S. Kounev. Online power consumption estimation for functions in cloud applications. In *2019 IEEE International Conference on Autonomic Computing (ICAC)*, pages 63–72. IEEE, 2019.
- [307] Y. R. Schneider and A. Koziolk. Towards reverse engineering for component-based systems with domain knowledge of the technologies used.
- [308] SEALABQualityGroup. E-shopper. <https://github.com/SEALABQualityGroup/E-Shopper>.
- [309] K. Seemakhupt, S. Liu, Y. Senevirathne, M. Shahbaz, and S. Khan. Pmnet: in-network data persistence. In *2021 ACM/IEEE 48th Annual International Symposium on Computer Architecture (ISCA)*, pages 804–817. IEEE, 2021.
- [310] K. Sellami, M. A. Saied, and A. Ouni. A hierarchical dbscan method for extracting microservices from monolithic applications. In *Proceedings of the International Conference on Evaluation and Assessment in Software Engineering 2022*, pages 201–210, 2022.
- [311] P. V. Seshadri, H. Balagopal, A. Nayak, A. P. Kumar, and P. Loyola. Conveyor move2kubernetes: A framework for automated application re-platforming. In *2022 IEEE 15th International Conference on Cloud Computing (CLOUD)*, pages 115–124. IEEE, 2022.
- [312] R. Sharma and A. Mathur. Traefik as kubernetes ingress. In *Traefik API Gateway for Microservices*, pages 191–245. Springer, 2021.
- [313] O. Shchur, A. C. Turkmen, T. Januschowski, J. Gasthaus, and S. Günemann. Detecting anomalous event sequences with temporal point processes. *Advances in Neural Information Processing Systems*, 34:13419–13431, 2021.
- [314] J. Shi, J. Wang, K. Fu, Q. Chen, D. Zeng, and M. Guo. Qos-awareness of microservices with excessive loads via inter-datacenter scheduling. In *2022 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*, pages 324–334. IEEE, 2022.
- [315] A. Singjai and U. Zdun. Conformance assessment of architectural design decisions on the mapping of domain model elements to apis and api endpoints. In *2022 IEEE 19th International Conference on Software Architecture Companion (ICSA-C)*, pages 76–83. IEEE, 2022.
- [316] A. Singjai, U. Zdun, O. Zimmermann, M. Stocker, and C. Pautasso. Patterns on designing api endpoint operations. 2021.
- [317] D. Sokolowski. Deployment coordination for cross-functional devops teams. In *Proceedings of the 29th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, pages 1630–1634, 2021.
- [318] D. Sokolowski. Infrastructure as code for dynamic deployments. 2022.
- [319] J. Soldani and A. Brogi. Automated generation of configurable cloud-native chaos testbeds. In *European Dependable Computing Conference*, pages 101–108. Springer, 2021.
- [320] J. Soldani, G. Montesano, and A. Brogi. What went wrong? explaining cascading failures in microservice-based applications. In *Symposium and Summer School on Service-Oriented Computing*, pages 133–153. Springer, 2021.
- [321] J. Soldani, G. Muntoni, D. Neri, and A. Brogi. The μ tosca toolchain: Mining, analyzing, and refactoring microservice-based architectures. *Software: Practice and Experience*, 51(7):1591–1621, 2021.
- [322] G. Somashekar, A. Dutt, R. Vaddavalli, S. B. Varanasi, and A. Gandhi. B-meg: Bottlenecked-microservices extraction using graph neural networks. 2022.
- [323] G. Somashekar and A. Gandhi. Towards optimal configuration of microservices. In *Proceedings of the 1st Workshop on Machine Learning and Systems*, pages 7–14, 2021.
- [324] H. Song, P. H. Nguyen, and F. Chauvel. Using microservices to customize multi-tenant saas: From intrusive to non-intrusive. In *Joint Post-proceedings of the First and Second International Conference on Microservices (Microservices 2017/2019)*. Schloss Dagstuhl-Leibniz-Zentrum für Informatik, 2020.
- [325] J. Sorgalla, P. Wizenty, F. Rademacher, S. Sachweh, and A. Zündorf. Applying model-driven engineering to stimulate the adoption of devops processes in small and medium-sized development organizations. *SN Computer Science*, 2(6):1–25, 2021.
- [326] S. Speth. Semi-automated cross-component issue management and impact analysis. In *2021 36th IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 1090–1094. IEEE, 2021.
- [327] A. Sriraman and T. F. Wenisch. μ suite: A benchmark suite for microservices. In *2018 IEEE International Symposium on Workload Characterization (IISWC)*, pages 1–12. IEEE, 2018.
- [328] A. Sriraman and T. F. Wenisch. μ tune: Auto-tuned threading for {OLDI} microservices. In *13th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 18)*, pages 177–194, 2018.
- [329] C.-a. Sun, J. Wang, Z. Liu, and Y. Han. A variability-enabling and model-driven approach to adaptive microservice-based systems. In *2021 IEEE 45th Annual Computers, Software, and Applications Conference (COMPSAC)*, pages 968–973. IEEE, 2021.
- [330] H. Sun, D. Bonetta, C. Humer, and W. Binder. Efficient dynamic analysis for node.js. In *Proceedings of the 27th International Conference on Compiler Construction*, pages 196–206, 2018.
- [331] Y. Sun, L. Zhao, Z. Wang, D. Cui, Y. Yang, and Z. Gao. Fault root rank algorithm based on random walk mechanism in fault knowledge graph. In *2021 IEEE International Symposium on Broadband Multimedia Systems and Broadcasting (BMSB)*, pages 1–6. IEEE, 2021.
- [332] Y. Tang, N. Chen, Y. Zhang, X. Yao, and S. Yu. Fine-grained diagnosis method for microservice faults based on hierarchical correlation analysis. In *CCF Conference on Computer Supported Cooperative Work and Social Computing*, pages 14–28. Springer, 2022.
- [333] P. Tennage, S. Perera, M. Jayasinghe, and S. Jayasena. An analysis of holistic tail latency behaviors of java microservices. In *2019 IEEE 21st International Conference on High Performance Computing and Communications; IEEE 17th International Conference on Smart City; IEEE 5th International Conference on Data Science and Systems (HPCC/SmartCity/DSS)*, pages 697–705. IEEE, 2019.
- [334] V. Thrivikraman, V. R. Dixit, V. K. Gowda, S. K. Vasudevan, and S. Kalambur. Misertrace: Kernel-level request tracing for microservice visibility. *CoRR*, 2022.
- [335] O. Tomarchio, D. Calcaterra, G. Di Modica, and P. Mazzaglia. Torch: a toska-based orchestrator of multi-cloud containerised applications. *Journal of Grid Computing*, 19(1):1–25, 2021.
- [336] K. A. Torkura, C. Meinel, and N. Kratzke. Don’t wait to be breached! creating asymmetric uncertainty of cloud applications via moving target defenses. *arXiv preprint arXiv:1901.04319*, 2019.
- [337] M. Toslali, E. Ates, A. Ellis, Z. Zhang, D. Huye, L. Liu, S. Puterman, A. K. Coskun, and R. R. Sambasivan. Automating instrumentation choices for performance problems in distributed applications with

- vaif. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 61–75, 2021.
- [338] M. Toslali, S. Parthasarathy, F. Oliveira, H. Huang, and A. K. Coskun. Iter8: Online experimentation in the cloud. In *Proceedings of the ACM Symposium on Cloud Computing*, pages 289–304, 2021.
- [339] L. Traini and V. Cortellessa. Delag: Detecting latency degradation patterns in service-based systems. *arXiv preprint arXiv:2110.11155*, 2021.
- [340] C. Trubiani, A. Bran, A. van Hoorn, A. Avritzer, and H. Knoche. Exploiting load testing and profiling for performance antipattern detection. *Information and Software Technology*, 95:329–345, 2018.
- [341] T. Ueda, T. Nakaïke, and M. Ohara. Workload characterization for microservices. In *2016 IEEE international symposium on workload characterization (IISWC)*, pages 1–10. IEEE, 2016.
- [342] T. N. Y. U. T. Ueda and M. Ohara. Performance characteristics of linux for java workloads oversubscribing memory.
- [343] A. van Hoorn and T. F. Düllmann. Continuous testing tool i. 2020.
- [344] W. Viktorsson, C. Klein, and J. Tordsson. Security-performance trade-offs of kubernetes container runtimes. In *2020 28th International Symposium on Modeling, Analysis, and Simulation of Computer and Telecommunication Systems (MASCOTS)*, pages 1–4. IEEE, 2020.
- [345] J. von Kistowski, S. Eismann, N. Schmitt, A. Bauer, J. Grohmann, and S. Kounev. TeaStore: A Micro-Service Reference Application for Benchmarking, Modeling and Resource Management Research. In *Proceedings of the 26th IEEE International Symposium on the Modelling, Analysis, and Simulation of Computer and Telecommunication Systems, MASCOTS '18*, September 2018.
- [346] H. Vural and M. Koyuncu. Does domain-driven design lead to finding the optimal modularity of a microservice? *IEEE Access*, 9:32721–32733, 2021.
- [347] A. Walker, D. Das, and T. Cerny. Automated microservice code-smell detection. In *Information Science and Applications*, pages 211–221. Springer, 2021.
- [348] L. Wang, N. Zhao, J. Chen, P. Li, W. Zhang, and K. Sui. Root-cause metric location for microservice systems via log anomaly detection. In *2020 IEEE International Conference on Web Services (ICWS)*, pages 142–150. IEEE, 2020.
- [349] Q. Wang, L. Shwartz, G. Y. Grabarnik, V. Arya, and K. Shanmugam. Detecting causal structure on cloud application microservices using granger causality models. In *2021 IEEE 14th International Conference on Cloud Computing (CLOUD)*, pages 558–565. IEEE, 2021.
- [350] R. Wang, G. Casale, and A. Filieri. Service demand distribution estimation for microservices using markovian arrival processes. In *International Conference on Quantitative Evaluation of Systems*, pages 310–328. Springer, 2021.
- [351] X. Wang, C. Li, L. Zhang, X. Hou, Q. Chen, and M. Guo. Exploring efficient microservice level parallelism. In *International Parallel and Distributed Processing Symposium (IPDPS)*, 2022.
- [352] S. Waterworth. Stan’s robot shop – a sample microservice application. <https://www.instana.com/blog/stans-robot-shop-sample-microservice-application/>, 2018.
- [353] M. Wawrzoniak, I. Müller, R. Bruno, A. Klimovic, and G. Alonso. Short-lived datacenter. *arXiv preprint arXiv:2202.06646*, 2022.
- [354] N. Wenzler. Automated performance regression detection in microservice architectures. B.S. thesis, 2017.
- [355] L. Wu, J. Bogatinovski, S. Nedelkoski, J. Tordsson, and O. Kao. Performance diagnosis in cloud microservices using deep learning. In *AIOPS 2020-International Workshop on Artificial Intelligence for IT Operations*, 2020.
- [356] L. Wu, J. Tordsson, A. Acker, and O. Kao. Microras: Automatic recovery in the absence of historical failure data for microservice systems. In *2020 IEEE/ACM 13th International Conference on Utility and Cloud Computing (UCC)*, pages 227–236. IEEE, 2020.
- [357] L. Wu, J. Tordsson, E. Elmroth, and O. Kao. Microrca: Root cause localization of performance issues in microservices. In *NOMS 2020-2020 IEEE/IFIP Network Operations and Management Symposium*, pages 1–9. IEEE, 2020.
- [358] X. Wu, S. Cheng, S. Yuan, and Z. Wang. A parallelism-based earliest finish time (pbeft) algorithm for workflow scheduling in clouds. In *2022 7th International Conference on Cloud Computing and Big Data Analytics (ICCCBDA)*, pages 22–26. IEEE, 2022.
- [359] Y. Xia. *Reducing the Length of Field-replay Based Load Testing*. PhD thesis, Concordia University, 2021.
- [360] R. Xin, P. Chen, and Z. Zhao. Causalrca: Causal inference based precise fine-grained root cause localization for microservice applications. *arXiv preprint arXiv:2209.02500*, 2022.
- [361] J. Xu, P. Chen, L. Yang, F. Meng, and P. Wang. Logdc: Problem diagnosis for declaratively-deployed cloud applications with log. In *2017 IEEE 14th International Conference on e-Business Engineering (ICEBE)*, pages 282–287. IEEE, 2017.
- [362] M. Xu and R. Buyya. Brownoutcon: A software system based on brownout and containers for energy-efficient cloud computing. *Journal of Systems and Software*, 155:91–103, 2019.
- [363] R. Yedida, R. Krishna, A. Kalia, T. Menzies, J. Xiao, and M. Vukovic. Lessons learned from hyper-parameter tuning for microservice candidate identification. In *2021 36th IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 1141–1145. IEEE, 2021.
- [364] R. Yedida, R. Krishna, A. Kalia, T. Menzies, J. Xiao, and M. Vukovic. Partitioning cloud-based microservices (via deep learning). *arXiv preprint arXiv:2109.14569*, 2021.
- [365] K. Yin and Q. Du. On representing resilience requirements of microservice architecture systems. *International Journal of Software Engineering and Knowledge Engineering*, 31(06):863–888, 2021.
- [366] G. Yu, P. Chen, H. Chen, Z. Guan, Z. Huang, L. Jing, T. Weng, X. Sun, and X. Li. Microrank: End-to-end latency issue localization with extended spectrum analysis in microservice environments. In *Proceedings of the Web Conference 2021*, pages 3087–3098, 2021.
- [367] G. Yu, P. Chen, and Z. Zheng. Microscaler: Automatic scaling for microservices with an online learning approach. In *2019 IEEE International Conference on Web Services (ICWS)*, pages 68–75. IEEE, 2019.
- [368] G. Yu, P. Chen, and Z. Zheng. Microscaler: Cost-effective scaling for microservice applications in the cloud with an online learning approach. *IEEE Transactions on Cloud Computing*, 2020.
- [369] U. Zdun, P.-J. Queval, G. Simhandl, R. Scandariato, S. Chakravarty, M. Jelic, and A. Jovanovic. Microservice security metrics for secure communication, identity management, and observability. *ACM Transactions on Software Engineering and Methodology*, 2022.
- [370] C. Zhang, X. Peng, C. Sha, K. Zhang, Z. Fu, X. Wu, Q. Lin, and D. Zhang. Deeptralog: Trace-log combined microservice anomaly detection through graph-based deep learning. 2022.
- [371] H. Zhang, A. Cardoza, P. B. Chen, S. Angel, and V. Liu. Fault-tolerant and transactional stateful serverless workflows. In *14th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 20)*, pages 1187–1204, 2020.
- [372] K. Zhang, A. Filieri, and R. Chatley. Kubedim: Self-adaptive service degradation of microservices-based systems. 2021.
- [373] L. Zhang, V. Anand, Z. Xie, Y. Vigfusson, and J. Mace. The benefit of hindsight: Tracing edge-cases in distributed systems. *arXiv preprint arXiv:2202.05769*, 2022.
- [374] Y. Zhang, Y. Gan, and C. Delimitrou. μ qsim: Enabling accurate and scalable simulation for interactive microservices. In *2019 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*, pages 212–222. IEEE, 2019.
- [375] Y. Zhang, W. Hua, Z. Zhou, E. Suh, and C. Delimitrou. Sinan: ML-based & qos-aware resource management for cloud microservices.

- [376] L. Zhao, Y. Yang, Y. Li, X. Zhou, and K. Li. Understanding, predicting and scheduling serverless workloads under partial interference. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 1–15, 2021.
- [377] L. Zhao, Y. Yang, K. Zhang, X. Zhou, T. Qiu, K. Li, and Y. Bao. Rhythm: component-distinguishable workload deployment in datacenters. In *Proceedings of the Fifteenth European Conference on Computer Systems*, pages 1–17, 2020.
- [378] X. Zhou, X. Peng, T. Xie, J. Sun, C. Ji, W. Li, and D. Ding. Fault analysis and debugging of microservice systems: Industrial survey, benchmark system, and empirical study. *IEEE Transactions on Software Engineering*, 2018.
- [379] X. Zhou, X. Peng, T. Xie, J. Sun, C. Ji, W. Li, and D. Ding. Delta debugging microservice systems with parallel optimization. *IEEE Transactions on Services Computing*, 2019.
- [380] X. Zhou, X. Peng, T. Xie, J. Sun, C. Ji, D. Liu, Q. Xiang, and C. He. Latent error prediction and fault localization for microservice applications by learning from system trace logs. In *Proceedings of the 2019 27th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, pages 683–694, 2019.
- [381] X. Zhou, X. Peng, T. Xie, J. Sun, W. Li, C. Ji, and D. Ding. Delta debugging microservice systems. In *2018 33rd IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 802–807. IEEE, 2018.
- [382] X. Zhou, X. Peng, T. Xie, J. Sun, C. Xu, C. Ji, and W. Zhao. Poster: Benchmarking microservice systems for software engineering research. In *2018 IEEE/ACM 40th International Conference on Software Engineering: Companion (ICSE-Companion)*, pages 323–324. IEEE, 2018.
- [383] L. Zhu. Decomposition tool i. 2019.
- [384] X. Zhu, G. She, B. Xue, Y. Zhang, Y. Zhang, X. K. Zou, X. Duan, P. He, A. Krishnamurthy, M. Lentz, et al. Dissecting service mesh overheads. *arXiv preprint arXiv:2207.00592*, 2022.
- [385] O. Zimmermann, D. Lübke, U. Zdun, C. Pautasso, and M. Stocker. Interface responsibility patterns: processing resources and operation responsibilities. In *Proceedings of the European Conference on Pattern Languages of Programs 2020*, pages 1–24, 2020.
- [386] O. Zimmermann, C. Pautasso, D. Lübke, U. Zdun, and M. Stocker. Data-oriented interface responsibility patterns: Types of information holder resources. In *Proceedings of the European Conference on Pattern Languages of Programs 2020*, pages 1–25, 2020.
- [387] O. Zimmermann, M. Stocker, D. Lübke, C. Pautasso, and U. Zdun. Introduction to microservice api patterns (map). 2019.
- [388] W. Zorgdrager, K. Psarakis, M. Fragkoulis, E. Visser, and A. Kat-sifodimos. Stateful entities: Object-oriented cloud applications as distributed dataflows. *arXiv preprint arXiv:2112.00710*, 2021.